

Pokretanje aplikacije

Aplikacija je testirana na sledeće načine

- local
- Spark Standalone (pseudo distribuiran)
- Spark Standalone i HDFS (pseudo distribuiran)
- Spark Standalone i HDFS na docker kontejnerima, korišćenjem BDE docker slika

Priprema aplikacije

Svi fajlovi potrebni za pokretanje aplikacije se mogu kreirati pokretanjem skripte `zip.sh` koja se nalazi u repozitorijumu. Ova skripta kreira folder `zip`, u koji kopira `main.py` i kreira paket `jobs.zip`. Ovaj fajl i paket se mogu direktno koristiti za `spark-submit`.

Argumenti aplikacije

`main.py --input input --output output --job {{ job }} [--debug]`

--input: ulazni direktorijum, podržava sve fajl-sisteme koje podržava Spark `file://`, `hdfs://`, itd.

--output: izlazni direktorijum, podržava sve fajlsisteme koje podržava Spark `file://`, `hdfs://`, itd. U tom direktorijumu se kreira poddirektorijum sa imenom `{{job_name}}_{{date}}`, u kojem se nalaze podaci

--debug: prikazuje međukorake (za manji skup podataka)

--job:

- `daily_statistics` ili `ds`
- `transportation_statistics` ili `ts`
- `ransportation_modes` ili `tm`

Pokretanje u lokalnom modu

```
$SPARK_HOME/bin/spark-submit --py-files jobs.zip --master local[4] main.py --input /home/ana/input --output /home/ana/output --job {{name}}
```


Pokretanje u Spark Standalone modu

Spark verzija 3.3.5. [Dokumentacija](#).

Potrebno je pokrenuti spark master i spark worker. Spark `start-all.sh` skripta podrazumevano pokreće master i jedan worker na localhost-u. Nakon toga se aplikacija može proslediti korišćenjem `spark-submit` skripte.

```
$SPARK_HOME/sbin/start-all.sh
$SPARK_HOME/bin/spark-submit --py-files jobs.zip --master spark://localhost:7077 --executor-cores {{cores}} --executor-memory {{memory}} main.py --input /home/ana/input --output /home/ana/output --job {{name}}
```

Primer izvršenja aplikacije:



3.5.0

Spark Master at spark://localhost:7077

URL: spark://localhost:7077

Alive Workers: 1

Cores in use: 4 Total, 0 Used

Memory in use: 1024.0 MB Total, 0.0 B Used

Resources in use:

Applications: 0 Running, 3 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (1)

Worker Id	Address	State	Cores	Memory	Resources
worker-20240125194205-192.168.0.15-37151	192.168.0.15:37151	ALIVE	4 (0 Used)	1024.0 MB (0.0 B Used)	

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

Completed Applications (3)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20240125200551-0002	TransportationStatisticsJob_24_01_25T20_05_50	4	1024.0 MB		2024/01/25 20:05:51	ana	FINISHED	3.2 min
app-20240125200007-0001	DailyStatisticsJob_24_01_25T20_00_06	4	1024.0 MB		2024/01/25 20:00:07	ana	FINISHED	3.2 min
app-20240125195813-0000	TransportationModesJob_24_01_25T19_58_12	4	1024.0 MB		2024/01/25 19:58:13	ana	FINISHED	12 s

Spark Standalone sa HDFS-om

1. Podešavanje HDFS-a u pseudo distribuiranom modu

Hadoop verzija 3.3.6 - [Apache Docs](#)

1. Podesiti java home
2. Instalirati ssh i pdsh, podesiti passphrasless ssh
3. Podesiti HDFS konfiguracione fajlove

```
$ export JAVA_HOME={{here}}
$ sudo apt-get install ssh
$ sudo apt-get install pdsh
$ ssh-keygen -t rsa -P '' -f ~/.ssh/id_rsa
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
$ chmod 0600 ~/.ssh/authorized_keys
```

\$HADOOP_HOME/etc/hadoop/core-site.xml:

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

\$HADOOP_HOME/etc/hadoop/hdfs-site.xml:

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
</configuration>
```

\$HADOOP_HOME/etc/hadoop.env.sh:

```
export PDSH_RCMD_TYPE=ssh
```

2. Pokretanje HDFS demona

Formatiranje namenode-a, pokretanje HDFS demona, kreiranje roditeljskog direktorijuma

```
$HADOOP_HOME/bin/hdfs namenode -format
$HADOOP_HOME/sbin/start-dfs.sh
$HADOOP_HOME/bin/hdfs dfs -mkdir -p /user/ana
```

Postavljanje podataka na HDFS:

```
$HADOOP_HOME/bin/hdfs dfs -put /home/ana/input input
```

3. Pokretanje Spark master-a

```
$SPARK_HOME/sbin/start-all.sh
```

4. Pokretanje aplikacije

```
$SPARK_HOME/bin/spark-submit --py-files jobs.zip --master spark://localhost:7077 --executor-cores {{cores}} --executor-memory {{memory}} main.py --input hdfs://localhost:9000/user/ana/input --output hdfs://localhost:9000/user/ana/output --job {{name}}
```

5. Dovlačenje podataka sa HDFS-a

```
$HADOOP_HOME/bin/hdfs dfs -get output /home/ana/output
```

6. Zaustavljanje Spark i HDFS daemona

```
$HADOOP_HOME/sbin/stop-dfs.sh  
$SPARK_HOME/sbin/stop-all.sh
```

Pokretanje na klasteru Docker container-a

Korišćena BDE verzija: 3.1.2-hadoop3.2 ([Dokumentacija](#))

Pokretanje klastera

Iz direktorijuma u kome se nalazi docker-compose fajl i .env fajl se može pokrenuti docker compose. Ovim se pokreće klaster docker kontejnera koji se sastoji od HDFS i spark demona.

```
$ docker network create bde --attachable  
$ docker compose up -d
```

Podaci se kopiraju u namenode kontejner. Pokreće se interaktivni terminal na namenode kontejneru, gde se mogu izvršavati HDFS komande. Kreira se roditeljski direktorijum i u njega se postavljaju iskopirani podaci.

```
$ docker cp input namenode:/input
$ docker exec -it namenode bash
# hdfs dfs -mkdir -p /user/ana
# hdfs dfs -put input /user/ana/input
```

Pokretanje kontejnera koji će biti povezan na istu docker mrežu kao klaster i pozivanje spark submit iz tog kontejnera:

```
$ docker run -it --network bde --env-file hadoop.env -p 4040:4040 --name spark bde2020/spark-base:3.1.2-hadoop3.2 bash

$ docker cp zip spark:/zip

# spark/bin/spark-submit --master spark://spark-master:7077 --py-files zip/jobs.zip
--executor-cores {{c}} --executor-memory {{m}} zip/main.py --input hdfs://namenode:9000/user/ana/input
--output hdfs://namenode:9000/user/ana/output --job {{name}}
```

Nakon uspešnog izvršenja, iskopirati podatak sa HDFSa na namenode, a zatim na host mašinu.


```
$ docker exec -it namenode bash
# hdfs dfs -get /user/ana/output output
$ docker cp namenode:/output /home/ana/output
```

Primeri izvršenja

Lista kontejnera koji se izvršavaju:

```
ana@ana-PC:~/PycharmProjects/GeoLife$ docker ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS      NAMES
PORTS
265f9a5a79fa   bde2020/spark-base:3.1.2-hadoop3.2 "bash"                 18 seconds ago Up 17 seconds spark
0.0.0.0:4040->4040/tcp, :::4040->4040/tcp
1f18f5200a0d   bde2020/spark-worker:3.1.2-hadoop3.2 "/bin/bash /worker.sh" 2 minutes ago  Up 2 minutes spark-worker-2
8081/tcp, 0.0.0.0:8072->8071/tcp, :::8072->8071/tcp
b2d821256979   bde2020/spark-worker:3.1.2-hadoop3.2 "/bin/bash /worker.sh" 2 minutes ago  Up 2 minutes spark-worker-1
0.0.0.0:8071->8071/tcp, :::8071->8071/tcp, 8081/tcp
5752d3e5a3e5   bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8 "/entrypoint.sh /run..." 2 minutes ago  Up 2 minutes (healthy) namenode
0.0.0.0:9000->9000/tcp, :::9000->9000/tcp, 0.0.0.0:9870->9870/tcp, :::9870->9870/tcp
f907784dbe50   bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8 "/entrypoint.sh /run..." 2 minutes ago  Up 2 minutes (healthy) datanode
9864/tcp
fb2a014cc750   bde2020/spark-master:3.1.2-hadoop3.2 "/bin/bash /master.sh" 2 minutes ago  Up 2 minutes spark-master
0.0.0.0:7077->7077/tcp, :::7077->7077/tcp, 6066/tcp, 8080/tcp, 0.0.0.0:8070->8070/tcp, :::8070->8070/tcp
```

Pristupanje web interfejsu spark mastera pokrenutog na docker kontejneru:

 **Spark Master at spark://fb2a014cc750:7077**

URL: spark://fb2a014cc750:7077

Alive Workers: 2

Cores in use: 16 Total, 0 Used

Memory in use: 13.3 GiB Total, 0.0 B Used

Resources in use:

Applications: 0 Running, 3 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (2)

Worker Id	Address	State	Cores	Memory	Resources
worker-20240126000907-172.19.0.5-42931	172.19.0.5:42931	ALIVE	8 (0 Used)	6.6 GiB (0.0 B Used)	
worker-20240126000907-172.19.0.6-41829	172.19.0.6:41829	ALIVE	8 (0 Used)	6.6 GiB (0.0 B Used)	

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

Completed Applications (3)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20240126001434-0002	TransportationModesJob_24_01_26T00_14_33	16	1024.0 MB		2024/01/26 00:14:34	root	FINISHED	26 s
app-20240126001403-0001	TransportationModesJob_24_01_26T00_14_02	16	1024.0 MB		2024/01/26 00:14:03	root	FINISHED	7 s
app-20240126001146-0000	TransportationModesJob_24_01_26T00_11_44	16	1024.0 MB		2024/01/26 00:11:46	root	FINISHED	7 s

Pristupanje web interfejsu namenode-a pokrenutog na docker kontejneru:

Hadoop

Overview

Datanodes

Datanode Volume Failures

Snapshot

Startup Progress

Utilities

Browse Directory

/user/ana/output/TransportationModes/job_24_01_26T00_14_33/transportation_modes

Go!

Show

25

entries

Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	0 B	Jan 26 01:15	3	128 MB	._SUCCESS	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	9 B	Jan 26 01:15	3	128 MB	part-00000-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	5 B	Jan 26 01:14	3	128 MB	part-00001-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	4 B	Jan 26 01:14	3	128 MB	part-00002-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	4 B	Jan 26 01:15	3	128 MB	part-00003-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	7 B	Jan 26 01:14	3	128 MB	part-00004-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	5 B	Jan 26 01:15	3	128 MB	part-00005-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	6 B	Jan 26 01:14	3	128 MB	part-00006-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	5 B	Jan 26 01:15	3	128 MB	part-00007-0ec5f544-7026-4c69-9b59-8d013e36433e-c000.csv	